

Enes Uzun

Senior Software Architect *GPU compute, LLM inference infrastructure*

uzunenes.com me@uzunenes.com linkedin.com/in/uzunenes github.com/uzunenes Istanbul, Türkiye · Remote

Summary

Senior software architect at Ford Otosan (Ford Motor Company joint venture) with 8 years in production ML. I run a 16-GPU Kubernetes fleet across three sites in two countries (H200, A100, L40S and V100) and the SGLang serving on it: about 25,000 requests a day from 250 developers' coding agents on two H200s, 93% of prompt tokens served from the prefix cache, p95 time to first token under 1 s. The same fleet runs LLM fine-tuning, Triton vision inference and an internal chatbot for 15,000 employees. I work on utilisation, tail latency and cost per token, and on finding out why a serving engine misbehaves in production. Before that I built real-time C/C++ vision systems that run in four car plants.

Experience

Senior Software Architect

Aug 2026 – Present

Ford Otosan (Ford Motor Company joint venture), Istanbul. Senior Software Engineer from Dec 2022 to Aug 2026.

- Run the inference service for 250 developers and their coding agents on two H200 GPUs (SGLang with speculative decoding): about 25,000 requests a day with ~96K-token prompts, 2.4B prompt tokens a day of which 93% are served from the prefix cache, and p95 time to first token under 1 s. Code never leaves the network and there are no per-token API costs.
- Run the Kubernetes GPU cluster across three sites in two countries (5 servers, 16 GPUs: 2× H200, 4× A100, 6× L40S, 4× V100) with Prometheus and Grafana monitoring. SGLang serving runs on the H200s; LLM fine-tuning and Triton vision inference run on the A100 and L40S nodes. It hosts the internal chatbot (15,000 users), the coding models and the production vision models.
- Diagnosed a production SGLang regression: after a client disconnected, requests kept decoding for up to 13.5 minutes and flooded the logs with tens of thousands of lines. Reproduced it on a separate GPU rig across SGLang v0.5.17, v0.5.19 and v0.5.20 to isolate the release that introduced it and the one that fixed it, and reported our production reproduction upstream ([sgl-project/sglang#36333](https://github.com/epicdataco/svglang/pull/36333)).
- Designed and built the company's LLM gateway. It decides per request whether the data may leave the company: confidential code and documents go to open models on our own GPUs, the rest to Azure OpenAI. Access goes through Entra ID (OAuth2/OIDC).
- Lead technical direction for a team of 9. The team's work received 2 global Ford and 2 local Koç Group innovation awards.

Software Engineer, Network & Systems

May 2019 – Dec 2022

Ford Otosan, Gölcük plant

- Built Gözcü, a no-code platform for real-time accident detection on CCTV, deployable to a new camera in one click.
- Wrote its inference backend in C/C++ (GStreamer, OpenCV, YOLO): 250+ concurrent RTSP streams, under 200 ms end-to-end latency.

Software Engineer

Jun 2018 – May 2019

Evstek Information Technologies, Istanbul

- Developed the image-analysis software of KE-BOT, a robotic system for FUE hair transplantation (C++, OpenCV). The work was published at IEEE.

Open source and projects

Ultralytics YOLO

May 2025

[PR #20553](#), merged in release 8.3.131

- Co-authored, with three others, the official YOLO11 example for Triton Inference Server in C++: gRPC client, FP16 preprocessing, NMS, CMake build.

Hugging Face TRL

Sep – Oct 2026

[Issue #7449](#); fix merged as [PR #7463](#)

- Found and root-caused a bug where prompt-completion training (SFT, DPO, KTO) on Gemma 4 12B/26B/31B put the loss on the wrong tokens and silently dropped short answers. Opened the issue and a patch; the maintainers merged a closely related fix that credits me as co-author, released in v1.14.2 ([write-up](#)).

Own projects

github.com/uzunenes

- [triton-server-hpa](#): scales Triton on Kubernetes by GPU utilisation (DCGM metrics, Prometheus adapter, HPA), with GPU time-slicing so it runs on one GPU.
- [k8s-ai-stack](#): self-hosted LLM stack on Kubernetes with vLLM, Ollama and OpenWebUI (OAuth2), plus throughput benchmarks.

Skills

- **Inference:** vLLM, SGLang, Triton Inference Server, TensorRT, OpenAI-compatible APIs
- **Platform:** Kubernetes, OpenShift, Docker, ArgoCD, NVIDIA DCGM, Prometheus, Grafana, MLflow, GitHub Actions, Azure DevOps
- **Languages:** Python, C, Bash (expert); C++, Go; Rust (reading and debugging)
- **ML and vision:** PyTorch, YOLO, RT-DETR, OpenCV, GStreamer
- **Systems:** Linux, strace, perf; TCP/UDP, gRPC, MQTT, RTSP
- **Spoken:** Turkish (native), English (professional)

Publications and talks

- *Multi-site Kubernetes for on-prem GPU workloads*. Talk at KCD Istanbul 2026 (Kubernetes Community Days, CNCF), Jul 2026.
- *Autoencoder-based video anomaly detection*. 2nd International Conference of Engineering Sciences (ICES), Dec 2023.
- Image analysis for the KE-BOT hair transplantation robot. IEEE Xplore, document [8719213](#).

Education

M.Sc. Electronics Engineering

2019 – Jan 2027 (expected)

Gebze Technical University

- Coursework completed 2019–2022; thesis started in 2026, graduation expected January 2027.
- Research on unsupervised video anomaly detection with autoencoders (ICES 2023 paper above).

B.Sc. Electronics and Telecommunication Engineering

2013 – 2018

Kocaeli University

- Thesis: real-time lane departure warning on an embedded platform.

Courses

Fast and Efficient LLM Inference with vLLM (DeepLearning.AI, 2026); Accelerate Your Model (Google DeepMind on DataCamp, 2026); Executive Leadership (Koç University, 2025); Certified Scrum Team Member (Scrum Inc., 2022).

Updated 9 Oct 2026